

Game-theoretic APT defense: Overview of potentially problematic R code & behavior

Based on:

R-Code implementation:

https://github.com/jku-lit-scsi/ComputersAndSecurity_RoboticsCaseStudies_Cut-The-Rope

Title: Game-theoretic APT defense: An experimental study on robotics

Presented by: Benjamin Jablonski

Institute: Johannes Kepler Universität Linz

Date: 26.02.2025

Part 1: weight_value = 1 in ctr-core_1.R

- In Ctr-Core we have:

```
attack_graph <- set_edge_attr(attack_graph,  
                             name = "weight",  
                             index = get.edge.ids(attack_graph, edgelist),  
                             value = 1)  # the virtual start is no obstacle towards the real entry point
```

- Which makes the claim that by setting the weight value of a new edge to 1 it makes it trivial to travel

Part 1: hardness_value = 1 in experiment_3.R & experiment_4.R

At the same time experiments 3 & 4 argue that a hardness value of 1 makes takes the edge trivial to travel

```
# Handle NA values
hardness[is.na(hardness)] <- 1
# fix missing hardness values: if we know nothing,
# we consider the edge easy (trivial) to traverse
```

Can both be true?

How does the weight value of a given edge relate to hardness?

Part 1: The Mir100 Graph

We can find the answer in the “attack_graph_MIR100.R” file

There we can see that the hardness values (here called “edgeProbs”) are manually stored in the attack graph as “edge_probabilities” which later will be extracted as “hardness” values.

```
edgeProbs <- c(0.111265, 0.111265, 0.47287625, 0.47287625, 0.47287625, 0.3449215,  
0.47287625, 1, 0.3449215, 0.47287625, 1, 1, 1, 0.47287625, 0.47287625, 0.47287625,  
0.47287625, 0.47287625, 0.47287625, 0.47287625, 0.3449215, 0.3449215, 0.3449215, 1)  
  
attack_graph <- set_edge_attr(attack_graph,  
                             name = "edge_probabilities",  
                             index = E(attack_graph),  
                             value = edgeProbs)
```

Later in experiments 3 & 4 file:

```
hardness <- edge_attr(attack_graph, "edge_probabilities", E(attack_graph, path=route))  
hardness[is.na(hardness)] <- 1 # fix missing hardness values: if we know nothing, we consider
```

Part 1: The Mir100 Graph - weight is defined via edgeProbs!

At the same time we clearly can see how the weight values have been calculated!

```
attack_graph <- set_edge_attr(attack_graph,  
                               name = "weight",  
                               index = E(attack_graph),  
                               value = -log(edgeProbs))
```

$\text{weight} = -\log(\text{edge_probability})$

Or conversely: $\text{edge_probability} = \exp(-\text{weight})$

```
attack_graph <- set_edge_attr(attack_graph,
                             name = "weight",
                             index = get.edge.ids(attack_graph, edgelist),
                             value = 1) # the virtual start is no obstacle toward
```

Conclusion:

- The way the graph has been defined we now know that
- $\text{hardness} = \text{edge_probability} = \exp(-\text{weight})$
- Thus setting the weight value to 1 in the ctr-core_1.R file was most likely an oversight
- Correct would have been to set it to 0 instead

Proposed fix:

- As long as we stick with this specific definition for the weights (weight = $-\log(\text{edge_probability})$), it only makes sense to set all default weight values to 0 for this graph because:
- Hardness = $\exp(-0) = 1$

Alternatively one could modify randomSteps to always calculate hardness from weights directly:

- hardness <- **exp**(-edge_attr(attack_graph, "weight", E(attack_graph, path=route)))

Part 2: Which hardness/weight value does R use for later calculations **if multiple parallel edges with different weight values are present?**

Part 2: Which hardness/weight value does R use for later calculations if multiple edges with different weight values are present?

Mir100 Graph has 3 such Attack paths:

```
Path 3: [0, 3, 6, 8, 'c(12,13,14,16)']
Parallel edges detected between 8 -> c(12,13,14,16):
Edge key=0, weight=0.0
Edge key=1, weight=0.7489220813074156
*** These parallel edges have DIFFERENT weights ***

Path 5: [0, 3, 8, 'c(12,13,14,16)']
Parallel edges detected between 8 -> c(12,13,14,16):
Edge key=0, weight=0.0
Edge key=1, weight=0.7489220813074156
*** These parallel edges have DIFFERENT weights ***

Path 8: [0, 2, 11, 'c(12,13,14,16)']
Parallel edges detected between 11 -> c(12,13,14,16):
Edge key=0, weight=1.064439873679208
Edge key=1, weight=0.7489220813074156
Edge key=2, weight=0.0
*** These parallel edges have DIFFERENT weights ***
```

- Since each of these paths has parallel edges with different weights....which of these does R use to calculate the entries in the Payoff matrix?

```
Payoff Matrix (probability of reaching target):
Row 1: 0.039406 0.171217 0.195410 0.136853 0.122740 0.106259 0.058041 0.122740 0.042335 0.092405 0.028743
Row 2: 0.000000 0.000000 0.195410 0.000000 0.122740 0.000000 0.058041 0.122740 0.042335 0.000000 0.028743
Row 3: 0.044279 0.118219 0.083333 0.136853 0.122740 0.106259 0.058041 0.122740 0.042335 0.092405 0.028743
Row 4: 0.044279 0.078813 0.000000 0.078813 0.000000 0.106259 0.058041 0.122740 0.042335 0.092405 0.028743
Row 5: 0.044279 0.039406 0.195410 0.039406 0.122740 0.039406 0.058041 0.122740 0.042335 0.039406 0.028743
Row 6: 0.044279 0.171217 0.195410 0.136853 0.122740 0.078813 0.000000 0.122740 0.042335 0.092405 0.028743
Row 7: 0.044279 0.171217 0.195410 0.136853 0.122740 0.106259 0.058041 0.000000 0.042335 0.092405 0.028743
Row 8: 0.044279 0.171217 0.195410 0.136853 0.122740 0.106259 0.058041 0.122740 0.000000 0.092405 0.028743
```


Part 2: Inconsistency in Experiment 3

- Original Payoff Matrix

Pay-off matrix in R:

Payoff Matrix (probability of reaching target):

Row 1:	0.039199	0.170457	0.091827	0.136217	0.057858	0.105744	0.057858	0.121826	0.042338	0.091950	0.028909
Row 2:	0.000553	0.002182	0.091827	0.001719	0.057858	0.001269	0.057858	0.121826	0.042338	0.001147	0.028909
Row 3:	0.044047	0.117656	0.039444	0.136217	0.057858	0.105744	0.057858	0.121826	0.042338	0.091950	0.028909
Row 4:	0.044047	0.078418	0.000518	0.078418	0.000518	0.105744	0.057858	0.121826	0.042338	0.091950	0.028909
Row 5:	0.044047	0.039199	0.091827	0.039199	0.057858	0.039199	0.057858	0.121826	0.042338	0.039199	0.028909
Row 6:	0.044047	0.170457	0.091827	0.136217	0.057858	0.078418	0.000518	0.121826	0.042338	0.091950	0.028909
Row 7:	0.044047	0.170457	0.091827	0.136217	0.057858	0.105744	0.057858	0.000518	0.042338	0.091950	0.028909
Row 8:	0.044047	0.170457	0.091827	0.136217	0.057858	0.105744	0.057858	0.121826	0.000518	0.091950	0.028909

- With hardness = $\exp(-\min_weight)$ we see a discrepancy for attack paths 3 & 5 (blue)

Pay-off Matrix in Python with: hardness = $\exp(-\min_weight)$

Payoff Matrix (probability of reaching target):

Row 1:	0.039406	0.171217	0.195410	0.136853	0.122740	0.106259	0.058041	0.122740	0.042335	0.092405	0.028743
Row 2:	0.000000	0.000000	0.195410	0.000000	0.122740	0.000000	0.058041	0.122740	0.042335	0.000000	0.028743
Row 3:	0.044279	0.118219	0.083333	0.136853	0.122740	0.106259	0.058041	0.122740	0.042335	0.092405	0.028743
Row 4:	0.044279	0.078813	0.000000	0.078813	0.000000	0.106259	0.058041	0.122740	0.042335	0.092405	0.028743
Row 5:	0.044279	0.039406	0.195410	0.039406	0.122740	0.039406	0.058041	0.122740	0.042335	0.039406	0.028743
Row 6:	0.044279	0.171217	0.195410	0.136853	0.122740	0.078813	0.000000	0.122740	0.042335	0.092405	0.028743
Row 7:	0.044279	0.171217	0.195410	0.136853	0.122740	0.106259	0.058041	0.000000	0.042335	0.092405	0.028743
Row 8:	0.044279	0.171217	0.195410	0.136853	0.122740	0.106259	0.058041	0.122740	0.000000	0.092405	0.028743

- With hardness = $\exp(-\max_weight)$ we see a discrepancy for attack path 8 (yellow)

Pay-off Matrix in Python with hardness = $\exp(-\max_weight)$

Payoff Matrix (probability of reaching target):

Row 1:	0.039406	0.171217	0.092405	0.136853	0.058041	0.106259	0.058041	0.042335	0.042335	0.092405	0.028743
Row 2:	0.000000	0.000000	0.092405	0.000000	0.058041	0.000000	0.058041	0.042335	0.042335	0.000000	0.028743
Row 3:	0.044279	0.118219	0.039406	0.136853	0.058041	0.106259	0.058041	0.042335	0.042335	0.092405	0.028743
Row 4:	0.044279	0.078813	0.000000	0.078813	0.000000	0.106259	0.058041	0.042335	0.042335	0.092405	0.028743
Row 5:	0.044279	0.039406	0.092405	0.039406	0.058041	0.039406	0.058041	0.042335	0.042335	0.039406	0.028743
Row 6:	0.044279	0.171217	0.092405	0.136853	0.058041	0.078813	0.000000	0.042335	0.042335	0.092405	0.028743
Row 7:	0.044279	0.171217	0.092405	0.136853	0.058041	0.106259	0.058041	0.000000	0.042335	0.092405	0.028743
Row 8:	0.044279	0.171217	0.092405	0.136853	0.058041	0.106259	0.058041	0.042335	0.000000	0.092405	0.028743

Part 2: Where does this discrepancy come from?

The main suspect: E(...) method in R

```
# steps determined by hardness to exploit
randomSteps <- function(route, attackRate = NULL, defenseRate = NULL) {
  hardness <- edge_attr(attack_graph, "edge_probabilities", E(attack_graph, path=route))
}
```

Documentation: When using E(...) with path parameter on graphs with parallel edges, it **arbitrarily** selects one edge instead of returning all matching edges, causing inconsistent behavior in the randomSteps function.

Source: <https://igraph.org/r/doc/E.html> (see under Path Argument)

Part 2: Conclusion & Proposed Solution

- Conclusion:
 - The $E(\dots)$ method can not be trusted for our graphs with parallel edges
- Solution
 - Since a lower weight translates to a higher probability of success (remember $\text{hardness} = \exp(-w)$) it makes sense that in the presence of multiple parallel edges the attacker would pick the **edge with the lowest weight**.
 - **Thus $\text{hardness} = \exp(-\text{min_weight})$ in python implementation**

```
# Yes hardness of 1 means path is trivial, hardness 0 means path is impossible
hardness.append(np.exp(-min_weight))
```